

## TODAY IN AI

- Barclays expands Anthropic partnership, deploying Claude AI to modernize software and improve client support; The Den uses ChatGPT Work to cut administrative time, accelerating growth; NVIDIA Blackwell GPUs power OpenAI's GPT6 Astra Ultrafast, delivering up to 8× faster token generation; OpenAI essays explore AI's role in routine work and research productivity limits; AWS releases opensource Strands Decider 2B, a lowcost decision model for agentic workflows; New arXiv paper proposes gradientaligned pair selection to enhance personalized preference optimization.

## TOP STORIES

### 1. Barclays scales Claude to upgrade operations and improve client experience

Anthropic | Models | importance 10/10 | 01 Oct

Barclays broadens its Anthropic deal, rolling out Claude across the bank to speed software development, modernize legacy systems, and boost operational efficiency, aiming for 50% developer adoption by 2026.

### 2. The Den frees up 10-15 hours a week to grow with ChatGPT Work

OpenAI | Models | importance 10/10 | 01 Oct

The Den social club adopts ChatGPT Work, saving 1015 weekly hours by automating grant and licensing applications, enabling faster expansion and focused decisionmaking.

### 3. How NVIDIA GPUs Help Accelerate OpenAI's GPT-6 Astra Ultrafast

NVIDIA Blog | Models | importance 9/10 | 01 Oct

OpenAI's GPT6 Astra Ultrafast, running on NVIDIA Blackwell GPUs, offers up to eightfold faster token generation, improving latency for coding agents and interactive apps.

### 4. The eternal complement

OpenAI | Tools | importance 9/10 | 01 Oct

An OpenAI essay argues that advanced AI will excel at repetitive tasks, highlighting limits in research productivity and the need for scalable execution capacity.

### 5. Amazon releases its own Jev clone as decision models flood the web

TechCrunch AI | Models | importance 8/10 | 01 Oct

AWS launches Strands Decider 2B, an opensource, lowlatency decision model built on Qwen3.52B, targeting workflow steps requiring calibrated choices.

### 6. Gradient-Aligned Pair Selection for Personalized Preference Optimization

arXiv cs.AI | Research | importance 7/10 | 02 Oct

A new arXiv paper introduces GradientAligned Pair Selection, a method to improve personalized preference optimization in AI systems.

## ALSO WORTH KNOWING

### Research

- Characterizing a Configuration Where Inference-Time PRM-Pruned Fragment Grafting Is Inert: Evidence from Three Reasoning LMs (arXiv cs.AI)

### Tools

- Google thinks SpaceX's Starship has to launch 1,800 times before space data centers get off the ground (TechCrunch AI)
- Opus 5.5 loves to tell you 'this matters' (and other AI writing tells) (TechCrunch AI)
- Introducing Olmo-core 3: Open, scalable training infrastructure for large MoEs (Hugging Face)
- Inside our months-long investigation into Kevin O'Leary's Utah data center debacle (The Verge AI)
- From Proposal to Verified Effect: Praxa, an Evidence-Bound Harness for Governed AI Agent Execution (arXiv cs.AI)

### Business

- How Albertsons Companies is reimagining retail from the inside out (OpenAI)
- OpenAI's new agent is a shot at Meta - but can it compete with free? (The Verge AI)

### Policy

- OpenAI cuts ties with 3 safety researchers, WSJ reports (TechCrunch AI)
- Judge dismisses antitrust lawsuits over Google's AI Overviews (The Verge AI)