

TODAY IN AI

- OpenAI releases 722 AI-generated math papers with formal proofs and compute details, sparking debate on research ethics.
- Jump Trading deploys GPT6 agents to automate complex quantitative research, highlighting AI's expanding role in finance.
- Anthropic expands its Cyber Verification Program, offering tiered access to powerful models for defensive and redteam security work.

TOP STORIES

1. Sharing AI progress in mathematics

OpenAI | Models | importance 10/10 | 06 Oct

OpenAI published 722 math manuscripts, including formal Lean proofs, compute estimates and reasoning summaries, via GitHub, following advisory group guidelines to improve transparency and citation practices.

2. How Jump Trading is scaling quant research with ChatGPT

OpenAI | Models | importance 10/10 | 06 Oct

Jump Trading uses GPT6 Astra agents to autonomously handle extensive quantitative research tasks, enabling recursive improvement and realtime steering while managing compliance risks.

3. Expanding the Cyber Verification Program

Anthropic | Tools | importance 9/10 | 06 Oct

Anthropic launches an expanded Cyber Verification Program with three access tiers, granting vetted security teams reduced-safeguard models for defensive and authorized redteam activities.

4. OpenAI just dropped 700 preprints of mathematical proofs and counterexamples

Hacker News | Research | importance 8/10 | 06 Oct

A Hacker News post highlights OpenAI's release of 700 preprints containing AI-generated mathematical proofs and counterexamples, drawing community attention.

5. OpenAI drops another batch of mathematical breakthroughs

The Verge AI | Research | importance 8/10 | 06 Oct

The Verge reports OpenAI's batch of 722 papers solves hundreds of open math problems, includes compute metrics, and raises concerns about ethical disclosure practices.

6. Utah to let AI examine patients and prescribe medication without human oversight

Hacker News | Policy | importance 8/10 | 06 Oct

Utah pilots an AI-only acne diagnosis and prescription service, initially limited and overseen by doctors, marking a first state-level experiment in autonomous medical AI.

ALSO WORTH KNOWING

Models

- Anthropic is giving startups a free year of Claude Team and \$1,000 in credits (TechCrunch AI)
- Google is about to remove free access to Gemini Flash and Pro (The Verge AI)
- EmbeddingGemma 2: an open, lightweight multimodal embedding model (Google DeepMind)
- Mistral Says Its New AI Model 'Le Chonk' Is the Best Open-Weight Offering Outside of China (Wired AI)

Research

- Mathematical manuscripts and supporting proof artifacts produced by OpenAI (Hacker News)

Tools

- Advancing computer use with Ironclad (OpenAI)
- How AI decision models could change content moderation (TechCrunch AI)

Business

- Atlassian and OpenAI expand partnership to turn enterprise knowledge into action (OpenAI)
- AI computing startup Lambda to raise \$4B ahead of planned IPO (TechCrunch AI)

Policy

- South Korea says AI agents appear to have been used to hack the country's banks (Hacker News)