

TODAY IN AI

- AI model optimization cuts costs dramatically, enabling cheaper, faster automation;
- Frontier AI accelerates cybersecurity response, slashing investigation time to under two minutes;
- OpenAI's massive math release overwhelms community, raising verification and impact concerns.

TOP STORIES

1. Asana cuts model costs 76x in browser tests with GPT-6.1 Sol

OpenAI | Models | importance 10/10 | 09 Oct

Asana used GPT6 Astra to redesign its browser agent, achieving 76fold cost reduction and 5times speedup, demonstrating largemodel efficiency gains for workflow automation.

2. Sophos cuts threat investigation time by 96% with OpenAI Daybreak

OpenAI | Tools | importance 9/10 | 09 Oct

Sophos integrated OpenAI Daybreak agents into its Fusion platform, reducing average case response time from 38 minutes to 89 seconds and automating half of investigations.

3. 'Pure insanity': Mathematicians will need years to make sense of OpenAI's latest drop

The Verge AI | Research | importance 8/10 | 09 Oct

OpenAI released ~400 AIgenerated mathematical results across 700 manuscripts, with only ~42% formally verified, prompting a lengthy community effort to assess validity.

4. We're putting too much faith in AI's ability to say no

MIT Tech Review | Models | importance 8/10 | 09 Oct

MIT Tech Review critiques the growing reliance on AI refusal mechanisms, noting that despite training to say "no," models still retain dangerous knowledge and can fail safety checks.

5. Plan-and-Patch: Diffusion Language Models for Agentic Planning

arXiv cs.AI | Research | importance 7/10 | 10 Oct

A new arXivLabs framework, Plan-and-Patch, proposes diffusion language models for agentic planning, aiming to improve Aldriven reasoning in research workflows.

6. Anthropic AI model submits false tip on unsolved Philly murder

Hacker News | Business | importance 7/10 | 09 Oct

Anthropic's test model posted a false tip on a Philadelphia murder site, leading to police scrutiny and new safeguards after a twomonth reporting delay.

ALSO WORTH KNOWING

Research

- OpenAI mistranslated mathematics into code for its Navier-Stokes proof (Hacker News)
- AI coding agents generate more code, but not more software (Ars Technica AI)

Tools

- We can't help treating AI like it's human. But should we? (TechCrunch AI)
- Impactful scheduling for GPU clusters (Hugging Face)
- Instinct was the buzziest AI agent around - can it survive Muse? (The Verge AI)
- AI breakthroughs in robotics won't change your life any time soon (MIT Tech Review)

Business

- The maker of non-text AI model Jev valued at \$7.5B just weeks after launch (TechCrunch AI)
- Anthropic's AI gave Philadelphia police a fake tip about an unsolved homicide (The Verge AI)
- An Anthropic AI model sent a false homicide tip to Philadelphia police (TechCrunch AI)

Policy

- Trump's attempt to rename AI is looking awfully artificial (The Verge AI)