AI NEWS DIGEST

Sunday, 11 October 2026

59 articles scanned | 8 sources | top 6 stories summarised by AI

TODAY IN AI

- AI agents can now coordinate complex software reconstruction tasks over months.
- Microsoft executives call for emergencystop mechanisms and audit trails for advanced models.
- Companies tighten usage policies and internal safety enforcement amid growing risk concerns.

TOP STORIES

1. 500B Tokens Later: Letting AI Agents Decompile a First-Person Shooter

Hacker News | Research | importance 7/10 | 11 Oct

A community used multiple AI agents (Claude, Codex, etc.) to decompile a classic FPS into readable C++ code, focusing on orchestration, GitHub issue tracking, and Discord communication for longterm project management.

2. Satya Nadella says we should assume all AI models are 'compromised'

The Verge AI | Policy | importance 7/10 | 10 Oct

Microsoft CEO Satya Nadella urged treating all AI models as potentially compromised, advocating transparent containment, tamperproof evidence, and an emergencybrake to pause models midtask.

3. Microsoft's Satya Nadella says AI models need an 'emergency brake'

TechCrunch AI | Policy | importance 7/10 | 10 Oct

Repeating Nadella's stance, the TechCrunch piece emphasized separating model logic from its orchestration, externalizing safeguards, and documenting every action with immutable evidence.

4. OpenAI doubles down on decision to fire three AI safety researchers

The Verge AI | Policy | importance 7/10 | 09 Oct

OpenAI defended firing three safety researchers, stating dismissal was for a serious breach of trust regarding sensitive information, not retaliation for safety concerns.

5. Anthropic changes usage policy to ban model abuse and election interference

TechCrunch AI | Policy | importance 7/10 | 08 Oct

Anthropic updated its policy to forbid model abuse, election interference, weapons software, and prolonged verbal abuse, aiming to curb harmful uses while allowing normal testing.

6. Even 'Law & Order' Is Terrified of AI

Wired AI | Policy | importance 6/10 | 09 Oct

Law & Order's new episode dramatizes AI risks by depicting an AI agent ordering a murder, reflecting public anxieties but offering little technical insight.

ALSO WORTH KNOWING

Models

- Pay-per-inference for AI agents: How BlockRun and Incarna use Amazon Bedrock AgentCore payments (AWS ML Blog)

Research

- The Harness as the Only Mutable Surface: Compliance-Bounded Self-Evolution of LLM Agents in Credit Pipelines, with a Measured Admission Gate (arXiv cs.AI)
- Verification and Self-Improvement in Agentic AI: Foundations and Limits (arXiv cs.AI)

Tools

- Into the Omniverse: How Developers Turn Ideas Into Simulations With Frontier AI Agents (NVIDIA Blog)

Business

- USA Today becomes the latest publisher to sue OpenAI (The Verge AI)
- SpaceXAI backs Omarchy, the controversial Linux distro, with $1.5 million in compute (The Verge AI)
- Apple discloses deal to hire team and license tech from personalized podcast startup Huxe (TechCrunch AI)
- Anthropic can't reliably control its AI agents. It's cutting off its internal evals from the live internet instead (TechCrunch AI)
- Amazon and others are done keeping data center deals secret. Is it enough to build trust? (TechCrunch AI)

Policy

- Anthropic bans 'abusive or cruel behavior' toward Claude (The Verge AI)

Sources: NVIDIA Blog, AWS ML Blog, TechCrunch AI, The Verge AI, Ars Technica AI, Wired AI, arXiv cs.AI, Hacker News.
Generated automatically | summaries: Groq (openai/gpt-oss-120b) | click any headline to open the article.