AI NEWS DIGEST

Tuesday, 10 November 2026

67 articles scanned | 11 sources | top 6 stories summarised by AI

TODAY IN AI

- AI agents are scaling software development productivity, as OpenAI's CodeRabbit claims 20x output gains.
- Anthropic's new Claude Fable and Mythos 5.2 models target harder technical tasks, with faster inference and lower costs.
- DeepMind demonstrates AIassisted formal verification of the Signal protocol, highlighting both promise and current limits.

TOP STORIES

1. CodeRabbit scales agentic software development and ships 20x more

OpenAI | Models | importance 10/10 | 09 Nov

OpenAI's CodeRabbit uses its models to manage AIgenerated code, enabling a single engineering team to ship software at roughly twenty times the previous rate, showcasing agentic development pipelines.

2. Control Room: A central system to track and stop model workloads across our research infrastructure

OpenAI | Policy | importance 9/10 | 09 Nov

OpenAI's Control Room provides a centralized dashboard to monitor model workloads across its research infrastructure and can halt runs instantly when safety or security concerns emerge.

3. Anthropic launches new Claude Fable and Mythos 5.2 models

The Verge AI | Models | importance 8/10 | 09 Nov

Anthropic released Claude Fable 5.2 and Claude Mythos 5.2, designed for demanding technical problems; the models claim better coding, reasoning, vision, faster output, and have solved several open scientific challenges during testing.

4. A New AI Tool From Google DeepMind Proved the Signal Protocol Is Secure

Wired AI | Research | importance 7/10 | 10 Nov

Google DeepMind used its AlphaProof Code framework to formally verify key cryptographic guarantees of the Signal protocol, confirming its security and illustrating AIaided formal verification's potential.

5. Anthropic reportedly plans to move forward with IPO plans

TechCrunch AI | Business | importance 7/10 | 10 Nov

Anthropic is proceeding with an IPO despite market volatility, aiming to set a valuation below earlier $1.5 trillion targets, while investors gauge the health of the broader AI boom.

6. Anthropic releases new Mythos and Fable models amid RSI concerns

TechCrunch AI | Models | importance 7/10 | 09 Nov

Anthropic's Mythos 5.2 and Fable 5.2 launch with improved coding, reasoning, vision, lower pricing, and faster inference; the release notes ongoing concerns about evaluating covert capabilities and recursive selfimprovement.

ALSO WORTH KNOWING

Models

- Claude Fable 5.2 launch offers (Hacker News)
- Claude Fable 5.2 and Mythos 5.2 (Hacker News)

Research

- Innovating with AI in the special chemicals industry (MIT Tech Review)
- Claude Fable 5.2 scores 65 on Artificial Analysis, highest so far (Hacker News)
- How CrossnoKaye grounds physical AI with Amazon Bedrock AgentCore (AWS ML Blog)

Tools

- Kuken Industries searches 70 years of drawings in minutes (OpenAI)
- Pluggable Agentic AI Framework: A seven-layer composable enterprise agentic AI reference architecture on Amazon Bedrock AgentCore (AWS ML Blog)

Business

- Your Experts Have Partners. There's HOPE for Keeping Them Together (Hugging Face)
- Jensen Huang's Troubles Leave the AI Industry Reeling (Wired AI)
- ChatGPT Enterprise now available in AWS Marketplace (OpenAI)

Sources: OpenAI, Hugging Face, NVIDIA Blog, AWS ML Blog, TechCrunch AI, The Verge AI, MIT Tech Review, Ars Technica AI, Wired AI, arXiv cs.AI, Hacker News.
Generated automatically | summaries: Groq (openai/gpt-oss-120b) | click any headline to open the article.