AI NEWS DIGEST

Friday, 13 November 2026

70 articles scanned | 11 sources | top 6 stories summarised by AI

TODAY IN AI

- Google launches Gemini 4 Hydrogen, a faster, cheaper model for complex search queries and agentic workflows.
- OpenAI's Codex boosts productivity for companies like Gondola and Holiday Extras, showing growing enterprise adoption.
- Hugging Face releases finetuning guide for 3D model TRELLIS.2, enabling taskspecific 3D generation.

TOP STORIES

1. Gondola doubles revenue as its team builds more with Codex

OpenAI | Business | importance 9/10 | 12 Nov

Gondola uses OpenAI Codex to increase its lead engineer's output by 50%, helping creatives gain credit and hire, while ChatGPT streamlines daily tasks, illustrating Codex's impact on productivity.

2. Holiday Extras is rethinking what it can build with Codex

OpenAI | Tools | importance 9/10 | 12 Nov

Holiday Extras leverages Codex to transform customer expertise into improved apps, higher booking conversion rates, and new growth opportunities, highlighting AI's role in enhancing travel services.

3. Google's new Gemini model is better at understanding complicated questions

The Verge AI | Models | importance 8/10 | 12 Nov

Google introduces Gemini 4 Hydrogen, claiming superior handling of multipart, typoladen, and specific queries with richer visual answers, now default in AI Mode across its services.

4. Google's latest Search upgrade promises better answers to your most complicated questions

TechCrunch AI | Models | importance 8/10 | 12 Nov

Google's Gemini 4 Hydrogen model improves Search by answering complex questions faster, offering better intent detection, visual results, and multistep task execution, positioned as a costeffective alternative to larger models.

5. Introducing Gemini 4 Hydrogen

Google DeepMind | Models | importance 8/10 | 12 Nov

Google DeepMind details Gemini 4 Hydrogen's speed and efficiency for everyday coding and agentic workflows, priced flat per token and integrated with Orchestrator for multimodel pipelines.

6. The Definitive Fine-Tuning Guide for 3D Generation Models

Hugging Face | Models | importance 8/10 | 12 Nov

Hugging Face provides scripts and LoRAbased finetuning for TRELLIS.2 3D generation, enabling developers to adapt the model to specific object categories with reduced compute overhead.

ALSO WORTH KNOWING

Models

- Edge0-145B: Frontier AI Capabilities on Your Own Computer (Hacker News)
- How Amplitude ships product improvements in 1-2 days with OpenAI (OpenAI)
- Google rolls out Gemini 4 Hydrogen, finally replacing 3.8 Flash (Ars Technica AI)
- Google Search gets a major upgrade with Gemini 4 Hydrogen (Google AI Blog)

Research

- AI chatbots are recommending fake retail websites, study finds (The Verge AI)

Tools

- Deploy frontier AI models on AWS Trainium with SGLang-Neuron (AWS ML Blog)
- Run agents securely and efficiently with the next generation AgentCore runtime (AWS ML Blog)
- The Jensen Huang scandal is hardly over (The Verge AI)
- More game developers are starting to use generative AI (The Verge AI)

Policy

- Supporting research on AI and children's safety (OpenAI)

Sources: OpenAI, Google DeepMind, Google AI Blog, Hugging Face, AWS ML Blog, TechCrunch AI, The Verge AI, MIT Tech Review, Ars Technica AI, Wired AI, Hacker News.
Generated automatically | summaries: Groq (openai/gpt-oss-120b) | click any headline to open the article.